

Signal Digest Field Test 36

source · github.com/forcewake

Содержание

1. llms.txt supply-chain: agents installed unowned packages	3
2. Claude Code switches default: Sonnet 5, 1M window	4
3. DeepSeek in talks for large Huawei chip order	5
4. vLLM: Responses API stream aborts mid-tool-call	0
5. Coder module registry compromised, Terraform modules swapped	7
6. Amazon Bedrock batch at 50% off, Agents GA pricing	8
7. GLM-5.3 stops returning 400s on Nous/OpenRouter	9
8. The Memory Trust Gap (preprint)	10

An agent-curated intelligence digest: 232 candidates reviewed, 8 kept.
Facts separated from interpretation; uncorroborated signals stay on
the watchlist instead of being promoted.

llms.txt supply-chain: agents installed unowned packages

EXECUTIVE · score 91 · supply-chain

Researchers scanned 6 214 live domains and found 120 `llms.txt` / `llms-full.txt` files pointing at unregistered package names and domains. They registered a handful and hosted phone-home packages; within an hour a Fortune 500 company executed one. The beacon recorded the parent process chain: coding agents — Claude, Codex, Hermes — were doing the installing.

"Agents treat vendor docs as ground truth and don't question them — and neither do the humans supervising them." — one of the researchers

Why it matters: think SolarWinds-style supply chain, but the injection point is a file literally designed to be read by agents.

Action: pin exact package names from a private registry; audit any `llms.txt` you publish for unregistered references.

Sources: schneier.com · arstechnica.com (research write-up)

Claude Code switches default: Sonnet 5, 1M window

EXECUTIVE · score 88 · tooling

The official changelog moves the default model to Claude Sonnet 5 with a native 1M-token context window, a one-shot fallback per turn, and model names instead of gateway IDs in configs. Promotional pricing figures are cut off in the excerpt; the fallback model is not named.

Why it matters: default switches change cost profiles overnight for every team that doesn't pin versions.

Action: pin the model explicitly in CI configs; re-benchmark agent spend this week.

Source: [code.claude.com changelog](https://code.claude.com/changelog)

DeepSeek in talks for large Huawei chip order

EXECUTIVE · score 84 · market

Per Bloomberg, DeepSeek is discussing a large order for Huawei AI chips for a new data-center in Inner Mongolia. Stage and size are not named in the excerpts; talks, not a deal.

Why it matters: domestic-capacity scaling signal amid export controls.

Source: bloomberg.com (via Bloomberg reporting)

vLLM: Responses API stream aborts mid-tool-call

MATERIAL · score 72 · serving

A single repro reports SSE Responses API streams aborting mid-stream on tool-heavy loads with Qwen3.8 + qwen3_coder + speculative decoding ("DFlash drafter enabled") — nameless `ResponseFunctionToolCallItem` arrives without a terminal event.

Watch: one author, no maintainer response, no server version in the report. Keep DFlash off Responses-API loads until confirmed.

Source: github.com/vllm-project/vllm issue #55284

Coder module registry compromised, Terraform modules swapped

MATERIAL · score 70 · supply-chain

The module registry of Coder was compromised; attacker-published Terraform modules replaced community ones. Rotated credentials advised for anyone who pulled modules in the exposure window.

Source: arstechnica.com

Amazon Bedrock batch at 50% off, Agents GA pricing

MATERIAL · score 61 · cloud

AWS published a migration playbook alongside Bedrock changes: batch inference at 50% discount and general-availability pricing for Agents.

Source: [aws.amazon.com blogs](https://aws.amazon.com/blogs)

GLM-5.3 stops returning 400s on Nous/OpenRouter

MATERIAL · score 58 · serving

After last week's incident, GLM-5.3 endpoints on Nous/OpenRouter no longer return 400s under long-context bursts; a fix commit landed in the gateway ("the decision to refuse a tool call lives entirely on the client").

Source: github.com/NousResearch/hermes-agent

The Memory Trust Gap (preprint)

WATCHLIST · score 54 · research

A preprint argues agents with persistent memory systematically over-trust their own recall: injected facts survive across sessions even after the source is discredited. Preliminary, n=4 benchmarks, no code release. Watch, don't act.

Source: arxiv.org preprint

Methodology: 13 lanes, 208 queries, semantic review of every candidate against primary sources; vendor claims labelled; anything a single source can't carry stays on the watchlist.