

Signal Digest Field Test 36

source · github.com/forcewake

An agent-curated intelligence digest: 232 candidates reviewed, 8 kept. Facts separated from interpretation; uncorroborated signals stay on the watchlist instead of being promoted.

Содержание

1. lms.txt supply-chain: agents installed unowned packages	2
2. Claude Code switches default: Sonnet 5, 1M window	3
3. DeepSeek in talks for large Huawei chip order	4
4. vLLM: Responses API stream aborts mid-tool-call	0
5. Coder module registry compromised, Terraform modules swapped	6
6. Amazon Bedrock batch at 50% off, Agents GA pricing	7
7. GLM-5.3 stops returning 400s on Nous/OpenRouter	8
8. The Memory Trust Gap (preprint)	9

llms.txt supply-chain: agents installed unowned packages

EXECUTIVE · score 91 · supply-chain

Researchers scanned 6 214 live domains and found 120 `llms.txt` / `llms-full.txt` files pointing at unregistered package names and domains. They registered a handful and hosted phone-home packages; within an hour a Fortune 500 company executed one. The beacon recorded the parent process chain: coding agents — Claude, Codex, Hermes — were doing the installing.

"Agents treat vendor docs as ground truth and don't question them — and neither do the humans supervising them." — one of the researchers

Why it matters: think SolarWinds-style supply chain, but the injection point is a file literally designed to be read by agents.

Action: pin exact package names from a private registry; audit any `llms.txt` you publish for unregistered references.

Sources: schneier.com · arstechnica.com (research write-up)

Claude Code switches default: Sonnet 5, 1M window

EXECUTIVE · score 88 · tooling

The official changelog moves the default model to Claude Sonnet 5 with a native 1M-token context window, a one-shot fallback per turn, and model names instead of gateway IDs in configs. Promotional pricing figures are cut off in the excerpt; the fallback model is not named.

Why it matters: default switches change cost profiles overnight for every team that doesn't pin versions.

Action: pin the model explicitly in CI configs; re-benchmark agent spend this week.

Source: [code.claude.com changelog](https://code.claude.com/changelog)

DeepSeek in talks for large Huawei chip order

EXECUTIVE · score 84 · market

Per Bloomberg, DeepSeek is discussing a large order for Huawei AI chips to equip a new data-center in Inner Mongolia. The stage of the talks and the order size are not named in the excerpts — this is negotiations reporting, not a signed deal, and single-source market stories stay capped below the top score band for exactly that reason.

Why it matters: if the order lands, it is a domestic-capacity scaling signal under export controls — training compute moving to sanctioned silicon at scale changes capacity planning for everyone downstream of Chinese model releases.

Watch: procurement volume, data-center power permits in Inner Mongolia, and any Huawei statement. Talks collapse more often than they convert.

Source: bloomberg.com (via Bloomberg reporting)

vLLM: Responses API stream aborts mid-tool-call

MATERIAL · score 72 · serving

A single repro reports SSE Responses API streams aborting mid-stream on tool-heavy loads with Qwen3.8 + qwen3_coder + speculative decoding ("DFlash drafter enabled") — nameless `ResponseFunctionToolCallItem` arrives without a terminal event.

Watch: one author, no maintainer response, no server version in the report. Keep DFlash off Responses-API loads until confirmed.

Source: github.com/vllm-project/vllm issue #55284

Coder module registry compromised, Terraform modules swapped

MATERIAL · score 70 · supply-chain

The module registry of Coder was compromised: attacker-published Terraform modules displaced community ones, meaning `terraform init` against the registry could pull attacker code into build pipelines with whatever credentials the plan runner holds. The exposure window is published; anyone who pulled modules in it should treat registry content as untrusted until pinned.

Why it matters: this is the second registry-style incident in one digest — the llms.txt research above is the same shape from the agent side. The pattern for the quarter is *trusted distribution points*, not zero-days: registries, docs files, package mirrors.

Action: pin module versions by digest; rotate credentials available to plan runners; audit what executed during the window.

Source: arstechnica.com

Amazon Bedrock batch at 50% off, Agents GA pricing

MATERIAL · score 61 · cloud

AWS published a migration playbook alongside two Bedrock changes: batch inference at a fifty percent discount and general-availability pricing for Agents. The playbook itself is the more useful artifact — it walks through request signing, model mapping and the guardrails config that most teams skip on the first pass.

Why it matters: batch pricing is the difference between nightly eval suites that run and nightly eval suites that get disabled when the bill arrives. If you are running evaluation or enrichment workloads request-by-request, this is the migration to make this quarter.

Action: move anything latency-tolerant to batch; re-check guardrail policies before switching, GA pricing changed per-token assumptions.

Source: [aws.amazon.com blogs](https://aws.amazon.com/blogs)

GLM-5.3 stops returning 400s on Nous/OpenRouter

MATERIAL · score 58 · serving

After last week's incident, GLM-5.3 endpoints on Nous/OpenRouter no longer return 400s under long-context bursts. The fix commit is small and telling: the decision to refuse a tool call now lives entirely on the client, where it belonged — the gateway had been rejecting requests it could not semantically judge.

Why it matters: if you route agent traffic through OpenRouter with fallbacks, silent gateway-side refusals masquerade as model failures and poison your retry logic. Worth knowing the failure class is gone before you remove workarounds.

Action: drop any 400-retry special-casing for this route; keep client-side refusal handling.

Source: github.com/NousResearch/hermes-agent

The Memory Trust Gap (preprint)

WATCHLIST · score 54 · research

A preprint argues that agents with persistent memory systematically over-trust their own recall: facts injected into memory during one session survive into later sessions even after the source is discredited, and explicit corrections decay faster than the original contamination. The mechanism the authors propose — provenance-weighted recall — is sensible; the evidence is four benchmarks and no code release.

Why it stays on the watchlist: n=4 with no reproduction is a prompt for your own testing, not a basis for architecture decisions. If you run long-lived agent memory, replicate the persistence test on your own stack this month; that is the cheap version of acting on it.

Source: arxiv.org preprint

Methodology: 13 lanes, 208 queries, semantic review of every candidate against primary sources; vendor claims labelled; anything a single source can't carry stays on the watchlist.